

Data Structures and Algorithms

Lecture 25: Balanced Binary Search Trees (AVL Trees)

Nopadon Juneam
Department of Computer Science
Kasetsart university

Outlines

- Tree Sort (BST sort) and its complexity
- Balanced and unbalanced binary search trees
- Height-balance property
- AVL trees
- Insertion on AVL trees

Trees Sort (BST Sort)

"out-of-place"

- *Tree Sort* is a not-in-place sorting algorithm. The algorithm is based on constructing a BST and perform an in-order traversal on the constructed BST to list a sorted output sequence.

- Tree Sort requires $\Theta(n)$ extra space to store the BST.

- The time complexity of Tree Sort is determined by the time taken to construct the BST given an input sequence.

- At worst, the complexity of the construction can be $O(n^2)$

- At best, the complexity of the construction can be $O(\underline{n \log n})$

Tree Sort :

1. Construct BST

2. Run in order $\Rightarrow \Theta(n)$

Phase 1: Construct BST $\Rightarrow T(n) = ?$
Phase 2: Traverse BST $\Rightarrow O(h)$

\leftarrow BST has n nodes

Case analysis


W.C. Complexity of Tree Sort

- **Proposition 1:** The worst case time complexity of Tree Sort is $O(n^2)$

W.C. After each insertion, the height of the BST increases by one.


Therefore, the complexity of constructing the BST is

$$\sum_{h=0}^{n-1} \Theta(h) = \Theta\left(\sum_{h=0}^{n-1} h\right) = \Theta(\underline{n^2})$$




Recall: Complexity of Operations on BST (1)



Operations	Complexity
Search	$O(h)$
Minimum	$O(h)$
Maximum	$O(h)$
Successor	$O(h)$
Predecessor	$O(h)$
Tree-Insert	$O(h)$
Tree-Delete	$O(h)$
<p><u>Remark:</u> h is the height of a binary tree.</p> <ul style="list-style-type: none"> - At worst, h can be $n-1$. W.C. - At best, h can be $\log_2(n+1)-1$. B.C. 	



Recall: Complexity of Operations on BST(2)



Binary Search Trees of Different Heights




unbalanced BST \approx BST of height $\Theta(n)$


Unbalanced Binary Search Trees

$$h = \underline{\underline{\Theta(\log n) \text{ in B.C.}}}$$

- **Unbalanced binary search tree:**
The height of the search tree is approximately linear in the number of nodes
 $h = \Theta(n)$
- If a binary search tree is unbalanced, the performance it achieves is no better than the linear data structures




1 (2) (3) (4) 5



Complete Binary Tree

- Ideally, we would like to have a binary search tree to be **complete**, that is, every level is filled
- The height of a complete binary tree is always $\log_2(n+1)-1$
- However, it is almost impossible to always maintain the search tree as a complete binary tree






AVL Property

$|height(L) - height(R)| \leq 1$ ^{simulation 1} \Rightarrow height of BST always $O(\log n)$

- **AVL property:** For every internal node v , the height of the subtrees rooted at children of v differs by at most 1
- **Node's height (subtree's height):** The height of a node u in a tree is recursively defined by:
 - If u is external, then the height of u is zero
 - Otherwise, the height of u is one plus the maximum height of a child of u




If u is NULL height of u is -1.



AVL Trees


- Any binary search tree that satisfies the AVL property is said to be an **AVL tree**
- An immediate consequence of the AVL property is that a subtree of an AVL tree itself is an AVL tree
- In other words, the difference between the heights of left and right subtrees cannot be more than one for all nodes
- AVL Trees are named after the initials of its inventors (Adelson, Velskii, and Landis 1962)



AVL Tree's Height Property

- **Proposition 2:** Let T be a binary tree that exhibits the AVL property. The height of the binary tree T is $O(\log n)$.

Pf. Let n_h be the **minimum** number of nodes in an AVL Tree of height h .
 $\leq 2 \cdot \log_2 n$



AVL Tree

$$n_h = 1 + n_{h-2} + n_{h-1}$$

1
internal
node
 \downarrow
 h
 \downarrow
 $h-1$

Boundary condition

$n_0 = 1$

$n_1 = 2$

$$n_h > n_{h-2} + n_{h-1} > n_{h-2} + n_{h-2} = 2n_{h-2}$$



Complexity of Search Operations on AVL Trees

Operations	Complexity
Search	$O(\log n)$
Minimum	$O(\log n)$
Maximum	$O(\log n)$
Successor	$O(\log n)$
Predecessor	$O(\log n)$

Insertion on AVL Trees

- The insert operation on an AVL tree T is the same as that on a regular binary search tree, but we must perform additional steps to make sure that T still has the AVL property
- **Step 1:** Perform insertion as in a regular binary search tree
- **Step 2:** Restructure T to restore the AVL property by performing rebalancing process through rotations

Insert 10, 20, 30, 40, 50



Subroutine: Left-Rotate

left rotate x as y $\frac{u}{\frac{v}{\frac{g}{\frac{1}{11}u}}$

- Left-Rotate(x, T): Left rotate subtree rooted at x so that x becomes the left child of y (the current right child of x) and the left child of y becomes the new right child of x .

Left-Rotate(x, T):

$y = x.\text{right}$

$B = y.\text{left}$

$y.\text{left} = x$

$x.\text{parent} = y$

$x.\text{right} = B$

$B.\text{parent} = x$

// Update heights of x & y

$x.\text{height} = 1 + \max(\text{height}(x.\text{left}), \text{height}(x.\text{right}))$

$y.\text{height} = 1 + \max(\text{height}(y.\text{left}), \text{height}(y.\text{right}))$


$\theta(1)$

height(u):

if ($u == \text{NULL}$):

return -1

return $u.\text{height}$



$O(1)$

inorder

$A \times B \times C$



$A \times B \times C$

Subroutine: Right-Rotate

- Right-Rotate(x, T): Right rotate subtree rooted at x so that x becomes the right child of y (the current left child of x) and the right child of y becomes the new left child of x .



```
Right-Rotate( $x, T$ ):  
   $y = x.left$   
   $B = y.right$   
   $y.right = x$   
   $x.parent = y$   
   $x.left = B$   
   $B.parent = x$   
  // Update heights of  $x$  &  $y$   
   $x.height = 1 + \max(\text{height}(x.left), \text{height}(x.right))$   
   $y.height = 1 + \max(\text{height}(y.left), \text{height}(y.right))$ 
```

```
height( $u$ ):  
  if ( $u == \text{NULL}$ ):  
    return -1  
  return  $u.height$ 
```




Rebalancing Process

- Suppose x is lowest node violating the height-balance property.
- Assume x is right-heavy (the left case is symmetric).
- **Right-right case:** If x 's right child is right-heavy or balanced, perform left rotation.
- **Right-left case:** Else, perform two rotations (right rotation, and left rotation).
- Then, continue up to x 's grandparent, great-grandparent, and so on.




Examples of Insertions on AVL Trees


Insert(23)




x = 29: left-left case




Done




Insert(55)



x=65: left-right case



Done



Operation: AVL-Insert

- AVL-Insert(z, u, T): insert a new node z for which $z.key = v$, $z.parent = \text{NULL}$, $z.left = \text{NULL}$, and $z.right = \text{NULL}$ into an appropriate position in the AVL subtree of T

AVL-Insert(z, u, T):

// Step 1: Perform an insertion as in regular BST

if ($u == \text{NULL}$):

return z

if ($z.key < u.key$):

$u.left = \text{AVL-Insert}(z, u.left, T)$

else:

$u.right = \text{AVL-Insert}(z, u.right, T)$

// Step 2: Update height of node u

$u.height = 1 + \max(\text{height}(u.left), \text{height}(u.right))$

$b = \text{checkBalanced}(u)$

// Step 3: If u becomes unbalanced, then there are 4 cases

// Right-right case

if ($b < -1$ and $z.key > u.right.key$):

Left-Rotate(u, T)

...

// Right-left case

if ($b < -1$ and $z.key > u.left.key$):

Right-Rotate($u.right, T$)

Left-Rotate(u, T)

return u

height(u):

if ($u == \text{NULL}$):

return -1

return $u.height$

node:

node left

node right

node parent

int key

int height

checkBalanced(u):

if ($u == \text{NULL}$):

return -1

return $\text{height}(u.left) - \text{height}(u.right)$

Complexity of Operations on AVL Trees

Operations	Complexity
Search	$O(\log n)$
Minimum	$O(\log n)$
Maximum	$O(\log n)$
Successor	$O(\log n)$
Predecessor	$O(\log n)$
AVL-insert	$O(\log n)$

W.C. Complexity of Tree Sort (based on AVL Tree)

- **Proposition 3:** The worst case time complexity of Tree Sort based on using AVL Tree is $O(n \log n)$

Construct BST + Inorder

$$O(n \log n) + O(n) = O(n \log n)$$

Implementation of AVL-Insert in C (1)

```
// C program to insert a node in AVL tree
#include<stdio.h>
#include<stdlib.h>

// An AVL tree node
struct node
{
    int key;
    struct node *left;
    struct node *right;
    struct node *parent;
    int height;
};

// A utility function to get maximum of two integers
int max(int a, int b)
{
    return (a > b)? a : b;
}

// A utility function to get the height of the tree
int height(struct node* node)
{
    if (node == NULL)
        return -1;
    return node->height;
}
```

Implementation of AVL-Insert in C (2)

// Helper function that allocates a new node with the given key

```
struct node* createNode(int key)
{
    struct node* node = (struct node*)
        malloc(sizeof(struct node));
    node->key = key;
    node->left = NULL;
    node->right = NULL;
    node->parent = NULL;
    node->height = 0; // new node is initially added at leaf
    return node;
}
```

// A utility function to left rotate subtree rooted with x

```
struct node* leftRotate(struct node* x)
{
    struct node* y = x->right;
    struct node* B = y->left;
    // Perform rotation
    y->left = x;
    x->parent = y;
    x->right = B;
    if(B != NULL)
        B->parent = x;
    // Update heights
    x->height = max(height(x->left), height(x->right))+1;
    y->height = max(height(y->left), height(y->right))+1;
    // Return new root
    return y;
}
```


Implementation of AVL-Insert in C (3)

// A utility function to right rotate subtree rooted at x

```
struct node* rightRotate(struct node* x)
{
    struct node* y = x->left;
    struct node* B = y->right;
    // Perform rotation
    y->right = x;
    x->parent = y;
    x->left = B;
    if(B != NULL)
        B->parent = x;
    // Update heights
    x->height = max(height(x->left), height(x->right))+1;
    y->height = max(height(y->left), height(y->right))+1;
    // Return new root
    return y;
}
```

// Get Balance factor of a node

```
int getBalance(struct node* node)
{
    if (node == NULL)
        return -1;
    return height(node->left)-height(node->right);
}
```

Implementation of AVL-Insert in C (4)

```
// Recursive function to insert a key in the subtree rooted with node and returns the new root of the subtree.
struct node* insert(struct node* node, int key)
{
    /* 1. Perform the normal BST insertion */
    if (node == NULL)
        return(createNode(key));
    if (key < node->key) {
        node->left = insert(node->left, key);
        node->left->parent = node;
    }
    else if (key > node->key) {
        node->right = insert(node->right, key);
        node->right->parent = node;
    }
    else // Equal keys are not allowed in BST
        return node;

    /* 2. Update height of this ancestor node */
    node->height = 1 + max(height(node->left), height(node->right));

    /* 3. Get the balance factor of this ancestor node to check whether this node became unbalanced */
    int balance = getBalance(node);
    // If this node becomes unbalanced, then there are 4 cases
    // Left Left Case
    if (balance > 1 && key < node->left->key)
        return rightRotate(node);
    // Right Right Case
    if (balance < -1 && key > node->right->key)
        return leftRotate(node);
    // Left Right Case
    if (balance > 1 && key > node->left->key)
    {
        node->left = leftRotate(node->left);
        return rightRotate(node);
    }
    // Right Left Case
    if (balance < -1 && key < node->right->key)
    {
        node->right = rightRotate(node->right);
        return leftRotate(node);
    }
    /* return the (unchanged) node pointer */
    return node;
}
```

Implementation of AVL-Insert in C (5)

// A utility function to print preorder traversal of the tree. The function also prints height of every node

```
void inorder(struct node* node)
{
    if(node != NULL)
    {
        inorder(node->left);
        printf("key: %d, height: %d\n", node->key, node->height);
        inorder(node->right);
    }
}
```

/* Driver program to test above function*/

```
int main()
{
    struct node *root = NULL;

    root = insert(root, 10);
    root = insert(root, 20);
    root = insert(root, 30);
    root = insert(root, 40);
    root = insert(root, 50);
    root = insert(root, 25);

    printf("Preorder traversal of the constructed AVL"
           " tree is \n");
    inorder(root);

    return 0;
}
```

```
nj@Nopadons-MacBook-Pro codes % ./AVL_tree
Inorder traversal of the constructed AVL tree is
key: 10, height: 0
key: 20, height: 1
key: 25, height: 0
key: 30, height: 2
key: 40, height: 1
key: 50, height: 0
nj@Nopadons-MacBook-Pro codes %
```